



Deep Learning

Lecture 8: Deep Learning Architectures (Pt II)

University of Agder,
Kristiansand, Norway

Prepared by: Hadi Ghauch^{*}, Hossein S. Ghadikolaei[†]

^{*} Telecom ParisTech

[†] Royal Institute of Technology, KTH

<https://sites.google.com/view/fundl/home>

April 2019


Outline

1. Convolutional Neural Networks

Recap: Deep Neural Networks (DNNs)

Multi-layer Perceptron:

- Output from some perceptrons are used in the inputs to other perceptrons




Feedforward Deep Neural Network (DNN):

- Multi-layer Perceptron with 'many' hidden layers
- Only forward connections: from input to hidden layer or output

Recap: Deep Neural Networks (DNNs)

DNN with J -layers: cascade of J filters + activation



In compact form: $z_j = \sigma_j(W_j z_{j-1} + b_j)$

$z_j \in \mathbb{R}^{d_j}$: output to layer j

$W_j \in \mathbb{R}^{d_j \times d_{j-1}}$: **Weights** for layer j (drop bias for simplicity)

$\sigma_j : \mathbb{R}^{d_j} \mapsto \mathbb{R}^{d_j}$: **Activation func** for layer j

Outline

1. Convolutional Neural Networks

Definition and Preliminaries

Recall: $\mathbf{w} \in \mathbb{R}^{d_1}$ is vector, $\mathbf{W} \in \mathbb{R}^{d_1 \times d_2}$ is matrix. Tensors generalize this definition

third-order tensor: $\overline{\mathbf{W}} \in \mathbb{R}^{d_1 \times d_2 \times d_3}$

vectors/matrices are special case of a third-order tensor
concatenation of d_3 matrices each of size $d_1 \times d_2$

Color RGB image with $d_1 \times d_2$ pixels: modeled as $d_1 \times d_2 \times 3$ tensor
 $d_1 \times d_2$ matrix for each channel/color, with 3 channel

CNNs: input to each layer and action of layer modeled by tensor
tensors of higher order also used in CNNs (skipped here)


Vectorization: $\text{vec}(\mathbf{W})$ stacks columns of \mathbf{W} one-by-one to form $d_1 d_2$ column vector

$\text{vec}(-)$ is one-to-one mapping (due to preset order of cols)
 $\text{vec}(\text{vec}(\overline{\mathbf{W}}))$ gives a $d_1 d_2 d_3$ column vector
tensors, matrices and vectors are equivalent via $\text{vec}(-)$

Convolutional Neural Networks

Convolutional Neural Networks (CNNs)

- variants of DNNs where at least one convolution layer used
- layer = **convolution**, **non-linear pooling**, or **fully connected**



Source: Trivedi, Kodor, 2017

Convolutional Neural Networks (CNNs)

Task: Classify a color image into C -classes

$$\overline{\mathbf{Z}}_1 \rightarrow \underbrace{\overline{\mathbf{W}}_1}_{\text{Layer 1}} \rightarrow \overline{\mathbf{Z}}_2 \rightarrow \underbrace{\overline{\mathbf{W}}_2}_{\text{Layer 2}} \rightarrow \overline{\mathbf{Z}}_3 \rightarrow \cdots \rightarrow \mathbf{z}_J$$

$\overline{\mathbf{Z}}_j \in \mathbb{R}^{d_1^j \times d_2^j \times d_3^j}$: input tensor of layer j

third order tensor of dimension $d_1^j \times d_2^j \times d_3^j$

each element indexed by (m^j, n^j, p^j) :

$m^j \in \{0, \dots, d_1^j\}$, $n^j \in \{0, \dots, d_2^j\}$, $p^j \in \{0, \dots, d_3^j\}$, $\forall j$

e.g., **color image**: (m^j, n^j) pixel position, p^j channel/color

$\overline{\mathbf{Z}}_1$: **input tensor to CNN** (colored image to classify)

$\mathbf{z}_J \in \mathbb{R}^C$: **CNN output**, a probability mass function (PMF).

$[z_J]_c$ probability that input belongs to class $c \in C$

$\overline{\mathbf{W}}_j$: **tensor models effect of layer $j \in [J]$**

each layer can be: **convolution**, **non-linear** or **pooling** layer

math expression depends on type of layer (next)

Convolution with a kernel

Ex: **Convolution** of a **matrix (2nd order tensor)**, \mathbf{A} , with a **kernel**, \mathbf{K}

$$\begin{bmatrix} 1 & 2 & 3 & 1 \\ 4 & 5 & 6 & 1 \\ 7 & 8 & 9 & 1 \end{bmatrix} \star \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} (1/4) = \begin{bmatrix} 3 & 4 & 11/4 \\ 6 & 7 & 17/4 \end{bmatrix} := \mathbf{B} \quad (1)$$

check board for calculations

Partition \mathbf{A} into **all possible** 2×2 **blocks**

Each block is **mapped into one entry** in \mathbf{B}

the average of that block in that example (reduce noise)

In general, kernels defined via convolutions for a given task:

denoising, edge detection, etc.

In this example convolution done with **2nd order tensor** (matrix) as kernel
e.g. convolution of grayscale image and one kernel matrix

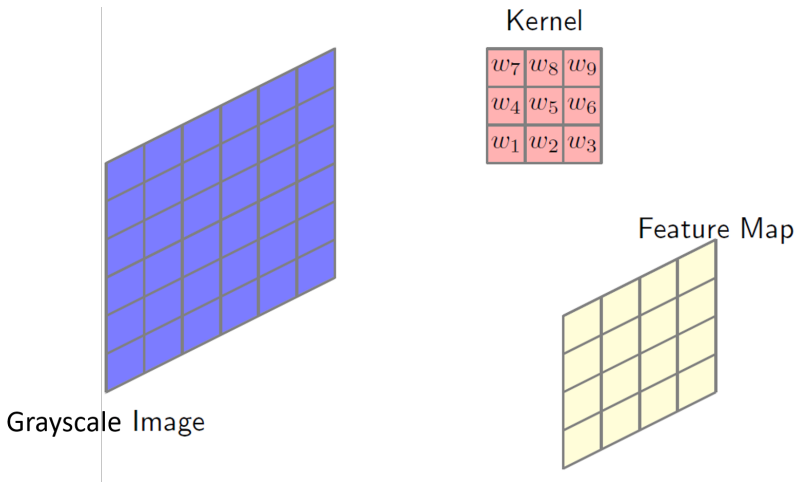
in practice convolution **color images** with **multiple kernel matrices**

convolution among two 3rd order tensors (potentially more)

(1) may be extended to convolutions of 3rd order tensors

Example


Convolution of grayscale image with Kernel matrix



Source: Trivedi, Kodor, 2017

Example


Convolution of grayscale image with Kernel matrix



Source: Trivedi, Kodor, 2017

Example


Convolution of grayscale image with Kernel matrix



Source: Trivedi, Kodor, 2017

Example

Convolution of grayscale image with Kernel matrix



Source: Trivedi, Kodor, 2017

Convolutional Layer

Layer j is **convolution** layer. Input-output relation :

$$\overline{\mathbf{Z}}_{j+1} = \overline{\mathbf{W}}_j \star \overline{\mathbf{Z}}_j$$

apply **more than one kernel matrix** \Rightarrow modeled as 3rd order tensor

$\overline{\mathbf{W}}_j \in \mathbb{R}^{D_1^j \times D_2^j \times D_3^j}$: set of **all kernels applied at layer j**

at layer j : **apply D_3^j kernels, each kernel is a $D_1^j \times D_2^j$ matrix**

$\overline{\mathbf{Z}}_j \in \mathbb{R}^{d_1^j \times d_2^j \times d_3^j}$: input tensor at layer j

$\overline{\mathbf{Z}}_{j+1} \in \mathbb{R}^{d_1^{j+1} \times d_2^{j+1} \times d_3^{j+1}}$: output tensor at layer j

convolution \star is among 3rd order tensors


expression generalization of 2D case in (1)

\star rewritten as matrix products


dimensions of $\overline{\mathbf{Z}}_{j+1}$ depends on that of $\overline{\mathbf{Z}}_j$ and $\overline{\mathbf{W}}_j$:

dimensions of $\overline{\mathbf{Z}}_{j+1}$ smaller if kernel larger 1×1


Convolutional Layer: Example




Convolutional Layer: Example



Convolutional Layer: Example



Convolutional Layer: Example



Nonlinear layer

Model for non-linear layers: $\overline{\mathbf{Z}}_{j+1} = \overline{\Sigma}_j(\overline{\mathbf{Z}}_j)$

$\overline{\Sigma}_j$: **non-linear element-by-element activation** for layer j

$\overline{\Sigma}_j : \mathbb{R}^{d_1^j \times d_2^j \times d_3^j} \mapsto \mathbb{R}^{d_1^{j+1} \times d_2^{j+1} \times d_3^{j+1}}$

$\overline{\Sigma}_j()$ element-by-element func: $d_1^{j+1} = d_1^j, d_2^{j+1} = d_2^j, d_3^{j+1} = d_3^j$

ReLU: commonly used for hidden layers

$$[\overline{\mathbf{Z}}_{j+1}]_{(m^j, n^j, p^j)} = \max(0, [\overline{\mathbf{Z}}_j]_{(m^j, n^j, p^j)})$$

$\forall m^j \in \{0, \dots, d_1^j\}, \forall n^j \in \{0, \dots, d_2^j\}, \forall p^j \in \{0, \dots, d_3^j\}$

non-linearity increases approximation power of model



Intuition: positive value implies presence of cat in image

negative value implies absence of cat in image

output of ReLU:


if cat absent (negative value), set output to 0

if cat present (positive value), pass to next layer



Softmax: used at output layer to ensure that \mathbf{z}_J is PMF
not an element-by-function

Nonlinear Layer: Example



Source: Trivedi, Kodor, 2017

Pooling layer

Pooling layer: $\bar{Z}_{j+1} = \bar{P}_j(\bar{Z}_j)$, $\bar{P}_j()$ pooling operator for layer j

$$\bar{P}_j : \mathbb{R}^{d_1^j \times d_2^j \times d_3^j} \mapsto \mathbb{R}^{d_1^{j+1} \times d_2^{j+1} \times d_3^{j+1}}$$

for each channel: **partition image into** $(D_1^j \times D_2^j)$ orthogonal blocks
combine/map each block into one scalar


depends only on (D_1^j, D_2^j) (no weights to learn)

dim of output tensor: $d_1^{j+1} = d_1^j / D_1^j$, $d_2^{j+1} = d_2^j / D_2^j$, $d_3^{j+1} = d_3^j$

subsampling and downconversion


Max pooling:

for each channel: partition image in
orthogonal $(D_1^j \times D_2^j)$ blocks of pixels
take max of each block



Average pooling:

for each channel: partition image in
orthogonal $(D_1^j \times D_2^j)$ blocks of pixels
take average of each block



Training a CNN

Training a CNN for image classification:

Training set: $\{\overline{\mathbf{X}}_i, \mathbf{y}_i\}_{i=1}^N$

$\overline{\mathbf{X}}_i$ 3rd order tensor representing (colored) image i (matrix for grayscale image)

\mathbf{y}_i vector has label of image i (cat, dog, tree, etc)

Loss function ?

ℓ_2 loss sometimes used, cross entropy loss for classification

Dersivations for backpropagation (chain rule) extended to 3rd order tensors;
see [Wu, 2017] section 6

Architectures and Variants

VGG-Verydeep-16:

pre-trained CNNs with 39 layers developed by Oxford, VGG

LeNet-5:

shallow CNN, few parameters to train (earlier arch)

AlexNet:

state-of-art arch for image classification (developed by A Krizhevsky, 2014)

Deep Scattering Transform [Mallat, 2014]

Deep CNNs use so-called scattering transforms

Convolution for each layer uses preset wavelet coeff. **No training!**

Many proofs showing its approximation power

Conclusions

Several DL architectures

- CNNs: mathematical model using tensors.
math model for each layer: convolutional, non-linear, pooling
some variants and practical implementations

Some references

J. Wu, "Introduction to convolutional neural networks ", May 1, 2017, available on arxiv

A. Goodfellow, Y. Bendgio, A Courville, "Deep Learning", MIT Press, Chap. 8, 9

Y. LeCun, L Bottou, G. Orr, K-B Muller, "Efficient Backprop", Neural Networks: Tricks of the trade

S. Trivedi, R. Kondor, CMSC 35246 Deep Learning, Spring 2017

C. M. Bishop, "Pattern Recognition and Machine Learning", Springer-Verlag, 2006

S. Haykin, "Neural Networks and Learning Machines", 3rd, 2009, Chap 4

Check Tutorial by Ruslan Salakhutdinov:
<http://www.cs.cmu.edu/~rsalakhu/>



Deep Learning

Lecture 8: Deep Learning Architectures (Pt II)

University of Agder,
Kristiansand, Norway

Prepared by: Hadi Ghauch^{*}, Hossein S. Ghadikolaei[†]

^{*} Telecom ParisTech

[†] Royal Institute of Technology, KTH

<https://sites.google.com/view/fundl/home>

April 2019